

Analysing customer reviews and comments with AI: 5 questions and prompts

For: marketing teams and founders of social enterprises **Data:** 358 product reviews from 8 platforms and 1,515 Instagram comments (all synthetic, English, for the fictional LoopCycle Nusantara) **Goal of the session:** show that AI can read hundreds of free-text comments, turn them into numbers, and help decide what to fix, say and measure, *and* that its output must be checked.

Suggested timing: 10 min intro, about 12 min per question, 5 min wrap-up.

1. What participants need

File	Rows	Contents
<code>loopcycle_reviews_amazon.csv</code> , <code>_shopee.csv</code> , <code>_tokopedia.csv</code> , <code>_tiktok_shop.csv</code> , <code>_lazada.csv</code> , <code>_blibli.csv</code> , <code>_shopify_website.csv</code> , <code>_google_store.csv</code>	358 in total	One file per platform. Column names differ by platform.
<code>loopcycle_instagram_comments.csv</code>	1,515	Comments on 8 Loop Tote posts, with replies
Optional: <code>loopcycle_tote_orders.csv</code>		Link reviews to orders (via the order ID column) for delivery-time follow-ups

Tell participants before they start:

1. The data is synthetic. Do not treat any figure as real.
2. The last two columns of the review and comment files, `synthetic_sentiment_label` and `synthetic_topic`, are **answer keys**. The prompts tell the AI not to use them until the checking step. For a stricter demo, delete those two columns first.
3. Platforms do not behave alike. Some files have no reply column, one has three ratings, and dates come in different formats.

2. Paste this first (shared context)

text

You are a customer-insight analyst helping a small social enterprise.

Context: LoopCycle Nusantara (Bandung, Indonesia) sells the Loop Tote, a tote b

Data (all English):

- Review files, one per platform: amazon, shopee, tokopedia, tiktok_shop, lazada
- Amazon: rating | review_title + review_text | no reply field | review_date |
- Shopee: rating_star | comment | seller_reply | create_time (Unix timestamp) |
- Tokopedia: rating | review_message | seller_response | review_time | invoice_
- TikTok Shop: review_rating | review_content | seller_reply | review_create_ti
- Lazada: product_rating (also seller_rating and delivery_rating) | review_titl
- Blibli: rating | review_content | merchant_reply | review_date (dd/mm/yyyy) |
- Shopify website: rating | review_title + review_body | merchant_reply | revie
- Google store: star_rating | review_text | owner_response | review_time | no o
- loopcycle_instagram_comments.csv: comments on 8 Loop Tote posts. Replies have
- The last columns named synthetic_sentiment_label and synthetic_topic are answ

Rules: use only the data I give you. Show how many reviews or comments sit behi

Warm-up demo (3 minutes)

- **Weak:** What do customers think of us?
- **Strong:** Question 1 below.

Compare the two answers. The weak prompt usually gives generic themes with no numbers, no checks and no next step.

3. The five questions

The path: **what do people say, is it the product, how do platforms differ, what stops people buying, what should we do.**

Question 1: What are customers really saying?

Why it matters: reading 358 reviews by hand takes hours. AI can do a first pass, but only if you verify it. **Files:** all 8 review files

text

Step 1: Combine the 8 review files into one table with these columns: platform,
Step 2: Using only the review text and rating (ignore the synthetic_ columns),
Step 3: Show the overall split of labels. Then show the top 5 topics for posi
Step 4: Now compare your labels with the synthetic_sentiment_label column. Show

Follow-ups:

- Which reviews should a human read first, and why?
- Where did the star rating mislead you, such as a positive review with 4 stars or an unrelated review with 4 stars?

Facilitator answer key (sample data)

- 358 reviews: Shopee 127, Shopify website 93, Tokopedia 46, Google store 33, Lazada 24, TikTok Shop 14, Amazon 13, Blibli 8.
- Answer-key split: **positive 208 (58%), unrelated 64 (18%), neutral 48 (13%), negative 38 (11%).**
- Top positive topics: quality, appearance, sustainability, usability. Top negative topic: quality (21 of the 38 negative reviews), then appearance (7). Top unrelated topics: delivery delay (21), packaging damage (14), seller service (10).
- **Expect the AI to do well** on clearly positive and clearly negative reviews. **Expect errors** between neutral and positive (3 and 4 stars), and on mixed reviews where the bag is fine but delivery was bad (6 reviews say the bag itself looked fine).
- **Rating traps:** 39 positive reviews are only 4 stars, 16 unrelated reviews are 4 stars, and 5 neutral reviews are 2 stars. Stars alone are not sentiment.
- **Wrangling traps:** Shopee dates are Unix timestamps, Blibli dates are dd/mm/yyyy, Lazada has three ratings (use `product_rating`), Amazon and the Google store file have no usable reply or order link.

Teaching point: AI does the first pass in minutes, but you must measure its accuracy against a trusted reference.

Question 2: Is the problem the product, or everything around the product?

Why it matters: star ratings blend product quality with logistics and service. Fixing the wrong thing wastes money. **Files:** all 8 review files (reuse the combined table from Question 1 if the chat is the same)

text

Is the problem the product, or everything around the product?

Use the combined review table (or rebuild it).

1. Split the reviews into: about the product, about delivery or packaging, about
2. Calculate the average rating overall and for product-focused reviews only. S
3. List the top product issues and the top non-product issues with counts and o
4. Show by platform how much of the low-rating problem is non-product. Flag pla
5. What does this mean for how we should read our star ratings, and which three

Follow-up: (If I also give you loopcycle_tote_orders.csv, how would you check whether delivery complaints go with longer delivery times? Describe the join and the columns first, then do it.)

Facilitator answer key

- Average rating **3.87**, but **4.12** when unrelated (non-product) reviews are removed.
- There are **69 one- and two-star reviews. 26 (38%) are not about the product**, 38 are product complaints and 5 are neutral.
- Product complaints: quality (21 of 38: stitching coming loose, plastic smell, rough fabric), appearance (7: colour faded or darker than photos), usability (4), value (3).
- Non-product complaints: delivery delay (21), packaging damage (14), seller service (10, plus 2 in store reviews), courier handling (6), wrong item (5), voucher or checkout trouble (4).
- By platform: Shopee 28 of 127 reviews and Tokopedia 10 of 46 are non-product (about 22% each). Amazon shows 5 of 13, but it is too small to rely on. Lazada's delivery rating averages 1.5 on the delivery-complaint reviews.
- Good fixes: stitching and smell checks before shipping, sturdier packaging and courier choice, and faster seller replies.

Teaching point: ask the AI to separate *what* is being complained about from *how unhappy* the customer is.

Question 3: How do platforms differ, and are we responding where it matters?

Why it matters: teams spread attention across platforms without knowing where customers are unhappy or ignored. **Files:** all 8 review files

text

How do platforms differ, and are we responding where it matters?

1. Compare the platforms: number of reviews, average rating, sentiment mix, ave
2. Show the seller reply rate by sentiment (positive, neutral, negative, unrela
3. Which kinds of reviews get no reply, and what is the risk or opportunity?
4. Identify two things that look like platform differences but could be caused
5. Recommend a reply policy: which reviews we answer within 24 hours, which lat

Follow-up: (Write three reply templates (delivery delay, product defect, neutral feedback) in our brand voice: warm, direct, no jargon.)

Facilitator answer key

- Platform picture: Shopee 127 reviews, average 3.79, 53% replied. Tokopedia 46, 3.59, 50%. Shopify website 93, 4.02, 40%, with the longest reviews (about 23 words against 12 to 13 on Shopee, Tokopedia and Google). Lazada 24, 3.92, 54%. TikTok Shop 14, 4.14, 21%. Blibli 8, 4.38, 88%. Amazon 13, 3.31, and **no reply field**. Google store 33, 4.09, 48%.
- **Reply rate by sentiment:** unrelated 85%, negative 84%, positive 36%, **neutral 24%**. Complaints get answered, neutral reviews rarely do.
- On Shopee, 29% of unrelated reviews include photos or video against 22% of positive and 6% of negative: photos are used as evidence for damaged or late parcels.
- **Traps:** Amazon, TikTok Shop and Blibli have 13, 14 and 8 reviews. Blibli's 88% reply rate rests on 8 reviews. A missing reply field on Amazon does not mean "no replies". Rating systems and review prompts differ by platform.

Teaching point: compare like with like, and flag when sample size makes a comparison meaningless.

Question 4: What are people saying on Instagram, and which comments stop them buying?

Why it matters: comments show pre-purchase questions and objections that reviews never capture. **Files:** (loopcycle_instagram_comments.csv) (optionally with

(loopcycle_post_kpis.csv) for post context)

text

What are people saying on Instagram, and which comments stop them buying?

Use `loopcycle_instagram_comments.csv`. Ignore the `synthetic_` columns until step

1. Separate top-level comments (`parent_comment_id` blank) from replies. Count ea
2. Classify top-level comments as: praise, question, objection (price, trust, q
3. For each group, calculate what share received a brand reply. Which groups do
4. Compare posts: which attract the most questions, objections and complaints?
5. Check your classification against `synthetic_topic` and `synthetic_sentiment_la`
6. Give 5 recommendations for managing comments, including three reusable repli

Follow-up: Estimate how many potential buyers these unanswered comments might represent, and say what we cannot know from this data.

Facilitator answer key

- **1,515 rows:** 1,130 top-level comments, 352 brand replies, 33 follower replies.
- Top-level sentiment: positive 509 (45%), neutral 389 (34%), unrelated 144 (13%), negative 88 (8%). Neutral is high partly because the quiz post attracted about 146 guesses.
- **The reply gap:** the brand replies to about **93% of delivery complaints** (26 of 28) and nearly all packaging and wrong-item complaints, but to only **2 of 24 price complaints (8%)**, **0 of 11 "is it really recycled?" doubts**, 4 of 13 quality complaints and 4 of 11 sold-out comments.
- Questions: 243 in total, 133 answered (55%). Where-to-buy 42 of 59, price 15 of 24, shipping 18 of 26.
- The batik drop post has the heaviest complaints (about 21% negative and 21% unrelated, mostly sold-out and checkout trouble), against about 4% negative on the launch Reel.
- Spam: 32 comments, 22 hidden. Another 27 are off-topic.
- **Traps:** quiz guesses ("14!") are neutral noise. Spam and off-topic comments count as "unrelated". Brand replies must be matched to parents through `parent_comment_id`.

Teaching point: the comments we ignore are often the persuasion barriers (price and trust), not the operational problems.

Question 5: What should we fix, say and measure next?

Why it matters: analysis only helps if it ends in decisions. **Files:** use the same chat as Questions 1 to 4, or attach all files.

text

Using the review and comment findings in this conversation (or rebuilding them

1. Compare themes across reviews and comments: what appears in both, only in re
2. Build a prioritised action list of at most 8 items across product, logistics
3. Write the content we should publish in response: three FAQ-style Instagram p
4. Propose 5 KPIs to track monthly (for example the share of 1 and 2 star revie
5. List what this data cannot tell us and what extra data we should collect.
6. Finally, argue against your own priorities: what would make them wrong?

Follow-ups:

- Turn this into a one-page brief for the founder.
- Which three actions could we start this week without spending money?

What good looks like

- Separates **pre-purchase friction** (price, trust, availability, shipping questions; mostly in comments) from **post-purchase problems** (quality, delivery, packaging; mostly in reviews).
- Spots the overlap: delivery and packaging issues show up in both sources.
- Prioritises: answering price and trust objections with content (the 14-bottle story, what the price includes, the QR traceability), stitching and smell checks, packaging and courier choices, a reply policy that includes neutral reviews and unanswered questions, and capacity for launches like the batik drop.
- Uses counts as evidence and labels small samples.
- Proposes KPIs that can actually be measured from these files.

Teaching point: ask the AI to critique itself, and keep a human responsible for the final choice.

4. Traps this session is built around

Trap	Where it shows up
Messy, mismatched files	Q1: 8 formats, different dates and rating columns
Stars are not sentiment	Q1, Q2: 4-star unrelated reviews, 2-star neutral reviews
Product vs everything else	Q2: 38% of low ratings are not about the bag
Tiny samples	Q3: Amazon (13), TikTok Shop (14), Blibli (8)
Noise that looks like signal	Q4: quiz guesses, spam, off-topic comments
What we ignore	Q4: price and trust objections go unanswered
Answer keys in the data	Q1, Q4: AI labels must be checked against a trusted reference

5. Keep participants honest

- **Spot-check:** each participant reads 10 reviews or comments and decides whether the AI's label is right.
- **Ask for the working:** (Show the rows behind this number.)
- **Ask for the opposite view:** (What would someone who disagrees say?)
- **Watch for over-confidence** on small samples, and for invented quotes. Check that every quote exists in the data.

6. Wrap-up question for the room

"Which finding would change what you do tomorrow, which would you check first before acting, and what would you want to ask customers directly?"